

Week8

September 20, 2023

Example :- Suppose the probability of the weather being cloudy is 40%. Also suppose the probability of rain on a given day is 20%. Also suppose the probability of clouds on a rainy day is 85%. If it's cloudy outside on a given day, what is the probability that it will rain that day?

```
[7]: def bayesTheorem(pA, pB, pBA):  
      return pA*pBA/pB  
  
pCloudy = 0.4  
pRain = 0.2  
pCloudyRain = 0.85  
  
bayesTheorem(pRain, pCloudy, pCloudyRain)
```

[7]: 0.425

1 Example

```
[9]: from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.naive_bayes import MultinomialNB  
from sklearn.metrics import accuracy_score, classification_report  
from sklearn.model_selection import train_test_split  
  
# Sample data  
texts = ["This is a positive review.", "Negative sentiment detected.", "A very_  
→positive experience.", "I didn't like this at all."]  
  
# Corresponding labels (1 for positive, 0 for negative)  
labels = [1, 0, 1, 0]  
  
# Split the data into a training set and a test set  
X_train, X_test, y_train, y_test = train_test_split(texts, labels, test_size=0.  
→2, random_state=42)  
  
vectorizer = CountVectorizer()  
X_train_vec = vectorizer.fit_transform(X_train)  
X_test_vec = vectorizer.transform(X_test)
```

```

clf = MultinomialNB()
clf.fit(X_train_vec, y_train)

y_pred = clf.predict(X_test_vec)

accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print(report)

```

Accuracy: 0.0

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1.0
1	0.00	0.00	0.00	0.0
accuracy			0.00	1.0
macro avg	0.00	0.00	0.00	1.0
weighted avg	0.00	0.00	0.00	1.0

```

/usr/lib/python3/dist-packages/sklearn/metrics/_classification.py:1221:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/usr/lib/python3/dist-packages/sklearn/metrics/_classification.py:1221:
UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0
in labels with no true samples. Use `zero_division` parameter to control this
behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

2 Questions

1. Implement in python of the following problems using Bayes Theorem.

- a) Of the students in the college, 60% of the students reside in the hostel and 40% of the students are day scholars. Previous year results report that 30% of all students who stay in the hostel scored A Grade and 20% of day scholars scored A grade. At the end of the year, one student is chosen at random and found that he/she has an A grade. What is the probability that the student is a hosteler?

```

[18]: pH, pD = 0.6, 0.4
      pAH, pAD = 0.3, 0.2

def bayes(h,d,ah,ad):
    return ((h*ah)/((h*ah)+(d*ad)))

```

```
prob = round((bayes(pH,pD,PAH,pAD)),4)
print(f"Probability of student being hosteller given A grade = {prob*100}%")
```

Probability of student being hosteller given A grade = 69.23%

- b) Suppose you're testing for a rare disease, and you have the following information:
- The disease has a prevalence of 0.01 (1% of the population has the disease).
 - The test is not perfect:
 - The test correctly identifies the disease (true positive) 99% of the time (sensitivity).
 - The test incorrectly indicates the disease (false positive) 2% of the time (1 - specificity).
- Calculate the probability of having the disease given a positive test result using Bayes' theorem.

```
[19]: pD, pN = 0.01, 0.99
      pTD, pTN = 0.99, 0.02

      def bayes(d,n,td,tn):
          num = d*td
          denom = (d*td)+(n*tn)
          return num/denom

      prob = bayes(pD,pN,pTD,pTN)
      print(f"Probability = {prob*100}%")
```

Probability = 33.33333333333333%

2. Write a program to implement the naïve Bayesian classifier without using scikit-learn library for the following sample training data set stored as a .CSV file. Calculate the accuracy, precision, and recall for your train/test data set. To classify 'If the weather is sunny, then the Player should play or not'?

```
[20]: import pandas as pd
      data = pd.read_csv('data.csv')
      data
```

```
[20]:      Outlook Play
0      Rainy   Yes
1      Sunny   Yes
2  Overcast   Yes
3  Overcast   Yes
4      Sunny   No
5      Rainy   Yes
6      Sunny   Yes
7  Overcast   Yes
8      Rainy   No
9      Sunny   No
10     Sunny   Yes
11     Rainy   No
12  Overcast   Yes
13  Overcast   Yes
```

```

[27]: class NaiveBayesClassifier:
    def __init__(self,X,y):
        self.X,self.y = X,y
        self.N = len(self.X) # Length of the training set
        self.dim = len(self.X[0]) # Dimension of the vector of features
        self.attrs = [[] for _ in range(self.dim)] # Here we'll store the
        ↪ columns of the training set

        self.output_dom = {} # Output classes with the number of occurrences in
        ↪ the training set. In this case we have only 2 classes

        self.data = [] # To store every row [Xi, yi]

        for i in range(len(self.X)):
            for j in range(self.dim):
                # if we have never seen this value for this attr before,
                # then we add it to the attrs array in the corresponding
                ↪ position
                if not self.X[i][j] in self.attrs[j]:
                    self.attrs[j].append(self.X[i][j])

                # if we have never seen this output class before,
                # then we add it to the output_dom and count one occurrence for now
                if not self.y[i] in self.output_dom.keys():
                    self.output_dom[self.y[i]] = 1
                # otherwise, we increment the occurrence of this output in the
                ↪ training set by 1
                else:
                    self.output_dom[self.y[i]] += 1
                # store the row
                self.data.append([self.X[i], self.y[i]])

        def classify(self, entry):
            solve = None # Final result
            max_arg = -1 # partial maximum

            for y in self.output_dom.keys():

                prob = self.output_dom[y]/self.N # P(y)

                for i in range(self.dim):
                    cases = [x for x in self.data if x[0][i] == entry[i] and x[1]
                    ↪ == y] # all rows with Xi = xi
                    n = len(cases)
                    prob *= n/self.N # P *= P(Xi = xi)

```

```

        # if we have a greater prob for this output than the partial
        ↪maximum...
        if prob > max_arg:
            max_arg = prob
            solve = y

    return solve

```

```

[28]: y = list(map(lambda v: 'Yes' if v == 1 else 'no', data['Play'].values)) #
        ↪target values as string
X = data[['Outlook']].values

y_train = y[:8]
y_val = y[8:]

X_train = X[:8]
X_val = X[8:]

nbc = NaiveBayesClassifier(X_train, y_train)

total_cases = len(y_val) # size of validation set

# Well classified examples and bad classified examples
good = 0
bad = 0

for i in range(total_cases):
    predict = nbc.classify(X_val[i])
    # print(y_val[i] + ' ----- ' + predict)
    if y_val[i] == predict:
        good += 1
    else:
        bad += 1

print('TOTAL EXAMPLES:', total_cases)
print('RIGHT:', good)
print('WRONG:', bad)
print('ACCURACY:', good/total_cases)

```

```

TOTAL EXAMPLES: 6
RIGHT: 6
WRONG: 0
ACCURACY: 1.0

```

```
[ ]:
```